



Rensselaer



# Autograding Distributed Algorithms in Networked Containers

Evan Maicus,  
Matthew Peveler,  
Stacy Patterson,  
Barbara Cutler



# Submittity

- A free, open source autograding platform
  - ~2500 users
  - 12-15 courses supported per term at RPI
  - In operation since 2014
- Support for:
  - Assignment Submission
  - Autograding
  - Exam Grading/Scanned PDF upload
  - Course Communications (Email/Forum)
  - Course Material Hosting
  - Plagiarism Detection



**New submission for: Homework 1**Due: 03/01/2019 @ 23:59

Drag your file(s) here or click to open file browser

By clicking Submittity you are confirming that you have read, understand, and agree to follow the Academic Integrity Policy.

Submit Clear Use Most Recent Submission

**Select Submission Version:** Version #1 Score: 5 / 20 GRADE THIS VERSION ▾ Do Not Grade This Assignment

*Note: This version of your assignment will be graded by the instructor/TAs and the score recorded in the gradebook.*

**Submitted Files**

buggy.cpp (0.08kb) 📎

submission timestamp: 02/24/2019 09:14:42 PM  
days late: 0 (before extensions)  
grading time: 4 seconds  
queue wait time: 0 seconds

**Results**

5 / 20 Total

5 / 5 Test 1 C++ - Compilation

0 / 15 Test 2 C++ - Execution [Details](#)


Student Program Output

Expected Program Output

```
1 Goodbye, world!
2
1 Hello, world!
2
```

[Visualize whitespace characters](#)

# Autograding Configuration



```
1 {
2   "testcases" : [
3     {
4       "type" : "Compilation",
5       "title" : "C++ - Compilation",
6       "command" : "clang++ -Wall -o a.out -- *.cpp",
7       "executable_name" : "a.out",
8       "points" : 5
9     },
10    {
11      "title" : "C++ - Execution",
12      "command" : "./a.out input_file.txt",
13      "points" : 15,
14      "validation" : [
15        {
16          "method" : "diff",
17          "actual_file" : "STDOUT.txt",
18          "description" : "Program Output",
19          "expected_file" : "test1_output.txt"
20        }
21      ]
22    }
23  ]
24 }
```

## New submission for: Homework 1

Due: 03/01/2019 @ 23:59

Drag your file(s) here or click to open file browser

By clicking Submit you are confirming that you have read, understand, and agree to follow the Academic Integrity Policy.

Submit

Clear

Use Most Recent Submission

Select Submission Version: Version #1 Score: 5 / 20 GRADE THIS VERSION Do Not Grade This Assignment

Note: This version of your assignment will be graded by the instructor/TAs and the score recorded in the gradebook.

### Submitted Files

buggy.cpp (0.08kb) 📎

submission timestamp: 02/24/2019 09:14:42 PM  
days late: 0 (before extensions)  
grading time: 4 seconds  
queue wait time: 0 seconds

### Results

5 / 20 Total

5 / 5 Test 1 C++ - Compilation

0 / 15 Test 2 C++ - Execution

[Details](#)

[Visualize whitespace characters](#)

#### Student Program Output


```
1 Goodbye, world!
2
```

#### Expected Program Output

```
1 Hello, world!
2
```

# Autograding Networked Assignments


- Course sizes are swelling, autograding can help [Wilcox, 2016]
- The number of students in advanced topics courses is swelling
- We aim to autograde networked assignments
- **Networked Assignment:** Two or more **instances** of software which communicate over a network.
- **Distributed Algorithm:** An algorithm carried out over connected **instances**.



# Complex Distributed Assignment

## Peer-To-Peer Itinerary

- A node can **schedule** an event
- **Accepted** if the time is available
- **Rejected** if the time is not available



# Challenges of Grading Distributed Assignments

## Logistics:

1. Prepare multiple instances
2. Network each test
3. Coordinate invocation and input
4. Isolate student nodes and networks

## Nondeterminism:

1. Issues of Concurrency
2. Network Stress
3. Issues of Distributed State  
[Beschastnikh et. al, 2015]

## Related Work:

1. Validate abstract postconditions [Marroquin et al. 2015]
2. Unit test system components [Torens et. al, 2010]
3. Inject network stress [Lubke et al. 2013, Alnawasreh et. al, 2017, PUMBA]

# Research Question

*Can we build a lightweight tool to help instructors assess student mastery of distributed assignments?  
Might this tool additionally provide benefits to student learning?*

# Outline: Requirements of Our System

- Provide powerful, simple network specification tools
- Facilitate scripted interaction
- Control network conditions
- Provide students with meaningful feedback



# Outline: Requirements of Our System

- Provide powerful, simple network specification tools
- Facilitate scripted interaction
- Control network conditions
- Provide students with meaningful feedback



## Docker Containers


- Easy to Configure
- Secure and Isolated
- Network tools

# Generating Reasonably Sized Student Networks

For each instance:

- **Host:** Where it's run
- **Connectivity:** Who it can talk to
- **Commands:** What it does
- **Environment:** What it has

```
1  "containers" : [  
2    {  
3      "container_name" : "server",  
4      "container_image" : "python:3.6",  
5      "outgoing_connections" : ["client"],  
6      "commands" : ["python3 server.py"]  
7    },  
8    {  
9      "container_name" : "client",  
10     "container_image" : "python:3.6",  
11     "outgoing_connections" : ["server"],  
12     "commands" : ["python3 client.py"]  
13   }  
14 ]
```



# Itinerary Specification

```
1 "containers" : [  
2   {  
3     "container_name" : "Peer A",  
4     "container_image" : "python:3.6",  
5     "commands" : ["python3 itinerary.py"]  
6   },  
7   {  
8     "container_name" : "Peer B",  
9     "container_image" : "python:3.6",  
10    "commands" : ["python3 itinerary.py"]  
11  },  
12  {  
13    "container_name" : "Peer C",  
14    "container_image" : "python:3.6",  
15    "commands" : ["python3 itinerary.py"]  
16  }  
17 ]
```

# Delivering Input to Student Applications

## Communication

- Standard Input


## Fault Tolerance

- Stop
- Start
- Kill

## Sequencing


- Delay

```
1  "input_actions" : [  
2    {  
3      "containers" : ["Client"],  
4      "action" : "stdin",  
5      "string" : "send:ping\n"  
6    }  
7  ]
```




# Itinerary Input

```
1  "input_actions" : [  
2    {  
3      "containers" : ["Peer A"],  
4      "action" : "stdin",  
5      "string" : "schedule B 8:00\n"  
6    },  
7    {  
8      "action" : "delay",  
9      "seconds" : 1  
10   },  
11   {  
12     "containers" : ["Peer C"],  
13     "action" : "stdin",  
14     "string" : "schedule B 8:00\n"  
15   },  
16   {  
17     "action" : "delay",  
18     "seconds" : 1  
19   },  
20   {  
21     "containers" : ["Peer A", "Peer B", "Peer C"],  
22     "action" : "stdin",  
23     "string" : "print itinerary\n"  
24   }  
25 ]
```




# Testing Fault Tolerance

```
1 "input_actions" : [  
2   {  
3     "containers" : ["Peer A"],  
4     "action" : "stdin",  
5     "string" : "schedule B 8:00\n"  
6   },  
7   {  
8     "action" : "delay",  
9     "seconds" : 1  
10  },  
11  {  
12    "containers" : ["Peer B"],  
13    "action" : "stop"  
14  },  
15  {  
16    "containers" : ["Peer C"],  
17    "action" : "stdin",  
18    "string" : "schedule B 8:00\n"  
19  },  
20  {  
21    "action" : "delay",  
22    "seconds" : 1  
23  },  
24  {  
25    "containers" : ["Peer A", "Peer B", "Peer C"],  
26    "action" : "stdin",  
27    "string" : "print itinerary\n"  
28  }  
29 ]
```



# Delivering Input to Student Applications

```
1 "input_actions" : [  
2   {  
3     "containers" : ["Peer A"],  
4     "action" : "stdin",  
5     "string" : "schedule B 8:00\n"  
6   },  
7   {  
8     "action" : "delay",  
9     "seconds" : 1  
10  },  
11  {  
12    "containers" : ["Peer B"],  
13    "action" : "stop"  
14  },  
15  {  
16    "containers" : ["Peer C"],  
17    "action" : "stdin",  
18    "string" : "schedule B 8:00\n"  
19  },  
20  {  
21    "action" : "delay",  
22    "seconds" : 1  
23  },  
24  {  
25    "containers" : ["Peer A", "Peer B", "Peer C"],  
26    "action" : "stdin",  
27    "string" : "print itinerary\n"  
28  }  
29 ]
```




# Providing Students with Meaningful Feedback

- It is difficult to determine a “happens before” relationship when debugging distributed algorithms [Beschastnikh et. al, 2015].

## The Logging Node

- Transparently injected into student applications.
- Intercepts all messages to provide a totally ordered message log
- **Network Aliasing**





# System Output

## C1 Log:

```
1   Initialized
2   Requested an appointment with C2 at 8:00
3   Appointment Scheduled
```

## C2 Log:


```
1   Initialized
2   C1 requested an appointment at 8:00
3   Accepting
4   C3 requested an appointment at 8:00
5   Rejecting
```

## C3 Log:

```
1   Initialized
2   Requested an appointment with C2 at 8:00
3   Appointment Conflict
```

## Logger Output:

```
1   Message passed from C1 to C2: "appt-8:00"
2   Message passed from C2 to C1: "accept"
3   Message passed from C3 to C2: "appt-8:00"
4   Message passed from C2 to C3: "conflict"
```



# Transparent Logging Node

## Advantages to the Student

- Provide an additional debugging resource

## Advantages to the Professor

- Impose network stress
- Enforces protocol use (TCP vs UDP)
- Restrict container communication
- Gauge program efficiency

## Limitations

- False Positive Message Receipt
- Sender Confusion
- Does not support Swarm Networks

# Case Study: Distributed Systems and Algorithms 2018

## Two Peer-To-Peer Assignments

- Team assignments
- Distributed Itinerary using log replication algorithm by Wu and Bernstein, then Paxos

## Live Demonstrations

- 20 minute face to face sessions with the instructor
- 7 cumulative testcases which increase in difficulty
- Adaptive expected output


## Our System:

- Used as a prerequisite for participation
- Post-semester study

# Case Study: Distributed Systems and Algorithms 2018

## Post-Semester Study

- 31 student pairs, 27 successfully graded
- Offline test with no feedback or resubmission
- 7 adapted testcases
- Pearson Correlation of .78
- Live demonstrations allow for more granular partial credit



# Autograding Complemented by Live Demonstration

## Advantages of Autograding:

- Multiple Submissions/early testing
- More testcase coverage
- Students can focus on the algorithm, rather than testing process
- Output is gathered to assist students in debugging

## Advantages of Live Demonstration:

- Better partial credit
- Adaptive corner case testing
- Dialog with the student

# Contributions


- **An Interface to Easily Create Student Networks**
  - Create networks and deploy student code
  - Automate timely interaction
  - Test program fault tolerance
- **Control Network Conditions with the Logging Node**
  - Simulate network stress
  - Enforce protocol use
- **Provide Useful Information to Students**
  - The logging node provides a totally ordered message log
- **Network generation is available as a standalone tool**

# Future Work

- **Additional Utilities for Controlling Network Conditions**
  - An interface to inject message delay
  - Chaos testing
- **Additional Output for Students**
  - Interleaved Output
- **Autograding Additional Advanced Topics Course**
  - Database Systems
  - Computer Graphics



# Rensselaer



Rensselaer Center for Open Source



<https://submittity.org/>

For More Submittity:

**Next Paper:**

Comparing Jailed Sandboxes  
vs Containers Within an  
Autograding System

**Tomorrow @3pm, Poster:**

Lichen: Customizable, Open Source  
Plagiarism Detection in Submittity

**Tomorrow @3pm, Poster:**

Facilitating Discussion-Based Grading  
And Private Channels via an  
Integrated Forum




# References

1. C. Wilcox, "Testing strategies for the automated grading of student programs," in Proceedings of the 47th ACM Technical Symposium on Computing Science Education, SIGCSE '16, (New York, NY, USA), pp. 437–442, ACM, 2016.
2. A. Marroquin, D. Gonzalez, and S. Maag, "Testing distributed systems with test cases dependencies architecture," in 2015 7th IEEE Latin-American Conference on Communications (LATINCOM), pp. 1–6, Nov 2015.
3. I. Beschastnikh, P. Wang, Y. Brun, and M. D. Ernst, "Debugging distributed systems," Queue, vol. 14, pp. 50:91–50:110, Mar. 2016.
4. R. Lübbe, D. Schuster, and A. Schill, "Reproducing network conditions for tests of large-scale distributed systems," in 2013 13th IEEE/ACM International Symposium on Cluster, Cloud, and Grid Computing, pp. 74–77, May 2013.
5. C. Torens and L. Ebrecht. 2010. RemoteTest: A Framework for Testing Distributed Systems.
6. K. Alnawasreh, P. Pelliccione, Z. Hao, M. Ränge, and A. Bertolino. 2017. Online Robustness Testing of Distributed Embedded Systems: An Industrial Approach. In 2017 IEEE/ACM 39th International Conference on Software Engineering: Software Engineering in Practice Track (ICSE-SEIP). 133–142.
7. TERRA NULLIS. 2018. Pumba - Chaos Testing for Docker. [https://alexei-led.github.io/post/pumba\\_docker\\_chaos\\_testing/](https://alexei-led.github.io/post/pumba_docker_chaos_testing/)

# Validating Distributed Algorithms

- **Instructor Client, Student Server.**
  - Instrumentation
- **Input Actions:** Test peer to peer. Test fault tolerance.
- **Validate Abstract Postconditions** [Marroquin et al. 2015]



# Capturing Output

## Peer A Log:


```
1   Initialized
2   Requested an appointment with Peer B at 8:00
3   Appointment Scheduled
```

## Peer B Log:


```
1   Initialized
2   Peer A requested an appointment at 8:00
3   Accepting
4   Peer C requested an appointment at 8:00
5   Rejecting
```

## Peer C Log:

```
1   Initialized
2   Requested an appointment with Peer B at 8:00
3   Appointment Conflict
```



# Using Docker Networks to Monitor and Restrict Communication



# Methods of Validation

- **Output Comparison Testing**

Student Program Output

```
1 #icebucketchallenge vs #alsicebucketchallenge, percentage change
2 150.0 vs 200.0
3 300.0 vs 400.0
4 500.0 vs 766.666666667
5 100.0 vs 92.3076923077
6 170.833333333 vs 320.0
7 7.69230769231 vs -19.0476190476
8
```

Expected Program Output

```
1 #icebucketchallenge vs #alsicebucketchallenge, percentage change
2 150 vs 200
3 300 vs 400
4 500 vs 766
5 100 vs 92
6 170 vs 320
7 7 vs -19
8
```

- **Script Based Testing:** Output is preprocessed using a Unix command


```
1 carrot      1 apple
2 apple      2 banana
3 banana      3 carrot
```

- **Output Analysis:** A custom algorithm is run on student output


```
1 1 + 2 = 3
2 3 - 4 = -1
3 5 + 4 = 9
```

# Limitation: False Positive Message Receipt

**Ideal:**



**Our System:**



# Complex Distributed Assignment

## Peer-To-Peer Itinerary

- A node can **schedule** an event
- **Accepted** if the time is available
- **Rejected** if the time is not available



# Docker Containers

Lightweight, secure environments virtualized at the operating system level and used to execute applications.

- Performant
- Easy to Configure
- Secure and Isolated
- Network tools





## Edit Gradeable: [hw01] Homework 1

For Help, Read: [Submitty Instructions on "Create or Edit a Gradeable"](#)

General

Submissions/Autograding

Rubric

Grader Assignment

Dates

Point precision:

Are student submissions/uploads a single file PDF with a fixed/prescribed format?

And do you expect each specific problem/item/component to appear on a specific page number of the PDF document?

☐ Yes ☒ No

**20 (0)** Total Points (Extra Credit Points)

**10**

Code Clarity



**10**

Adherence to the Project Specification

**✖Cancel**

Title

Note to TAs

Note to Students

Points

Additional Extra Credit Points  *Student scores will be clamped from above to Points + Additional Extra Credit Points*

Negative Penalty Points  *Student scores will be clamped from below to min(0, Negative Penalty Points)*

How should points be assigned? ☒ Grade by Count Up (from zero) ☐ Grade by Count Down (from "Points")

No Credit



Full Credit

Show mark to all students: ☐

Add New Mark

Add New Component

Export Components

Import Components

# Student Submission/Autograding

[Submitty](#) > [sample](#) > [Open Homework](#)

- ★ Gradeables
- 🔔 Notifications
- 📁 New Gradeable
- ⚙️ Course Settings
- 📖 Course Materials
- 💬 Discussion Forum
- 👤 Manage Students
- 📊 Manage Graders
- 👥 Manage Sections
- 📷 Student Photos
- 📅 Late Days Allowed
- 📅 Excused Absence Extensions
- ⚠️ Plagiarism Detection
- 📊 Grade Reports
- 📅 My Late Days/Extensions
- ☰ Collapse Sidebar
- 🔑 Logout Quinn

## New submission for: Open Homework

Due: 12/31/9996 @ 23:59

☒ Normal Submission ☐ Make Submission for a Student ☐ Bulk Upload

Drag your file(s) here or click to open file browser

By clicking Submit you are confirming that you have read, understand, and agree to follow the Academic Integrity Policy.

Submit

Clear

Use Most Recent Submission

Select Submission Version: Version #3 GRADE THIS VERSION ▼

Do Not Grade This Assignment

Note: This version of your assignment will be graded by the instructor/TAs and the score recorded in the gradebook.

### Submitted Files

part1\_buggy2.py (0.48kb) 📎

submission timestamp: 02/08/2019 05:30:09 PM  
days late: 0 (before extensions)  
grading time: 2 seconds  
queue wait time: 0 seconds

### Results

Test 1 Lab 1 Checkpoint 1 python \*.py

[Details](#)

Visualize whitespace characters

#### Student Program Output 📄

```
1 #icebucketchallenge vs #alsicebucketchallenge, percentage change
2 150.0 vs 200.0
3 300.0 vs 400.0
4 500.0 vs 766.666666667
5 100.0 vs 92.3076923077
6 170.833333333 vs 320.0
7 7.69230769231 vs -19.0476190476
8
```

#### Expected Program Output 📄

```
1 #icebucketchallenge vs #alsicebucketchallenge, percentage change
2 150 vs 200
3 300 vs 400
4 500 vs 766
5 100 vs 92
6 170 vs 320
7 7 vs -19
8
```